



Nodes of the

Cluster

[About](#)
[Nodes](#)
[Node Labels](#)
[Applications](#)
[NEW](#)
[NEW SAVING](#)
[SUBMITTED](#)
[ACCEPTED](#)
[RUNNING](#)
[FINISHED](#)
[FAILED](#)
[KILLED](#)
[Scheduler](#)

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Used Resources
0	0	0	0	0	<memory:0 B, vCores:0>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes
1	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Capacity
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers
/default-rack	RUNNING	hadoop-2274802010480:32835	hadoop-2274802010480:8042	Fri Mar 28 16:25:56 +0700 2025		0	

Showing 1 to 1 of 1 entries

← → ↺ 🏠 ⓘ http://localhost:9871/dfshealth.html#tab-overview

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Started:	Fri Mar 28 16:38:00 +0700 2025
Version:	3.4.0, rbd8b77f398f626bb7791783192ee7a5dfaec760
Compiled:	Mon Mar 04 13:35:00 +0700 2024 by root from (HEAD detached at release-3.4.0-RC3)
Cluster ID:	CID-f3690cda-0a60-4603-b957-5bf5affd7357
Block Pool ID:	BP-608357847-172.23.0.76-1743154673555

Copy_to_local:

```
from snakebite.client import Client

client = Client('localhost', 9000)
for f in client.copyToLocal(['/input/input.txt'], '/tmp'):
    print f
```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_han
h/HadoopWithPython/python/HDFS$ python2 copy_to_local.py
{'path': '/tmp/input.txt', 'source_path': '/input/input.txt', '
result': True, 'error': ''}
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_han
h/HadoopWithPython/python/HDFS$ cat /tmp/input.txt
jack be nimble
jack be quick
jack jumped over the candlestickhadoop@hadoop-2274802010480:~/i
h/HadoopWithPython/python/HDFS$ h
```

Delete.py

```
from snakebite.client import Client

client = Client('localhost', 9000)
for p in client.delete(['/foo', '/input'], recurse=True):
    print p
```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_han
h/HadoopWithPython/python/HDFS$ python2 delete.py
{'path': '/foo', 'result': True}
{'path': '/input', 'result': True}
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_han
h/HadoopWithPython/python/HDFS$
```

Get-pip.py

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$ python2 get-pip.py
DEPRECATION: Python 2.7 reached the end of its life on January 1st, 2020. Please upgrade your Python as Python 2.7 is no longer maintained. pip 21.0 will drop support for Python 2.7 in January 2021. More details about Python 2 support in pip can be found at https://pip.pypa.io/en/latest/development/release-process/#python-2-support pip 21.0 will remove support for this functionality.
Defaulting to user installation because normal site-packages is not writeable
Collecting pip<21.0
  Using cached pip-20.3.4-py2.py3-none-any.whl (1.5 MB)
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 20.3.4
    Uninstalling pip-20.3.4:
      Successfully uninstalled pip-20.3.4
Successfully installed pip-20.3.4
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$

```

List_directory.py

```

from snakebite.client import Client

client = Client('localhost', 9000)
for x in client.ls(['/']):
    print x

```

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$ python2 list_directory.py
{'group': u'supergroup', 'permission': 493, 'file_type': 'd', 'access_time': 0L, 'block_replication': 0, 'modification_time': 1743155609645L, 'length': 0L, 'blocksize': 0L, 'owner': u'hadoop', 'path': '/input'}
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$

```

List_file.py:

```

from snakebite.client import Client

client = Client('localhost', 9000)
for x in client.ls(['/input.txt']):
    print x

```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$ python2 list_file.py
{'group': u'supergroup', 'permission': 420, 'file_type': 'f', 'access_time': 1743156439018L, 'block_replication': 1, 'modification_time': 1743156439233L, 'length': 61L, 'blocksize': 134217728L, 'owner': u'hadoop', 'path': '/input.txt'}
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$
```

Mkdir.py:

```
from snakebite.client import Client

client = Client('localhost', 9000)
for p in client.mkdir(['/foo/bar', '/input'], create_parent=True):
    print p
```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$ python2 mkdir.py
{'path': '/foo/bar', 'result': True}
{'path': '/input', 'result': False, 'error': "mkdir: `/input': File exists"}
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$
```

Text.py:

```
from snakebite.client import Client

client = Client('localhost', 9000)
for l in client.text(['/input/input.txt']):
    print l
```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$ python2 text.py
jack be nimble
jack be quick
jack jumped over the candlestick
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/HDFS$
```

LAB2 Luigi:

Wordcount.py:

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ python2 wordcount.py WordCount
--input-file input.txt --output-file wordcount_output.txt --local-scheduler
DEBUG: Checking if WordCount(input_file=input.txt, output_file=wordcount_output.txt) is complete
DEBUG: Checking if InputFile(input_file=input.txt) is complete
INFO: Informed scheduler that task WordCount_input_txt_wordcount_output_0a6198752c has status PENDING
INFO: Informed scheduler that task InputFile_input_txt_fd27c12e07 has status DONE
INFO: Done scheduling tasks
INFO: Running Worker with 1 processes

```

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ cat wordcount_output.txt
be      2
jumped  1
over    1
candlestick    1
nimble  1
jack    3
quick   1
the     1
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$

```

Luigi_pig.py:

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ python2 -m luigi --module luigi_pig WordCount --input-file /user/hadoop/input.txt --output-file /user/hadoop/output --scheduler-host localhost --scheduler-port 8084
DEBUG: Checking if WordCount(input_file=/user/hadoop/input.txt, output_file=/user/hadoop/output, script_path=/home/hadoop/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/pig/wordcount.pig) is complete
DEBUG: Running file existence check: hadoop fs -stat /user/hadoop/output
INFO: Informed scheduler that task WordCount__user_hadoop_inp__user_hadoop_out__home_hadoop_iDr_10ee23b60a has status DONE
INFO: Done scheduling tasks
INFO: Running Worker with 1 processes
DEBUG: Asking scheduler for work...
DEBUG: Done
DEBUG: There are no more tasks to run at this time
INFO: Worker Worker(salt=695430859, workers=1, host=hadoop-2274802010480, username=hadoop, pid=20892) was stopped. Shutting down Keep-Alive thread
INFO:
===== Luigi Execution Summary =====

```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ hadoop fs -ls /user/hadoop/output
Found 2 items
-rw-r--r-- 1 hadoop supergroup      0 2025-03-29 00:34 /user/hadoop/output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    65 2025-03-29 00:34 /user/hadoop/output/part-r-00000
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ hadoop fs -cat /user/hadoop/output/part-r-00000
2      be
1      the
3      jack
1      over
1      quick
1      jumped
1      nimble
1      candlestick
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$
```

Luigi_mapreduce.py:

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ python2 luigi_mapreduce.py --input-file hdfs:///user/hadoop/input.txt --output-file hdfs:///user/hadoop/output/ --local-scheduler
No handlers could be found for logger "luigi-interface"
DEBUG: Checking if WordCount(input_file=hdfs:///user/hadoop/input.txt, output_file=hdfs:///user/hadoop/output/) is complete
DEBUG: Running file existence check: hadoop fs -stat hdfs:///user/hadoop/output/
INFO: Informed scheduler that task WordCount_hdfs__user_had_hdfs__user_had_c64bc50fc7 has status DONE
INFO: Done scheduling tasks
INFO: Running Worker with 1 processes
DEBUG: Asking scheduler for work...
DEBUG: Done
DEBUG: There are no more tasks to run at this time
INFO: Worker Worker(salt=587456706, workers=1, host=hadoop-2274802010480, username=hadoop, pid=36491) was stopped. Shutting down Keep-Alive thread
INFO:
===== Luigi Execution Summary =====

Scheduled 1 tasks of which:
* 1 complete ones were encountered:
  - 1 WordCount(input_file=hdfs:///user/hadoop/input.txt, output_file=hdfs:///user/hadoop/output/)

Did not run any tasks
This progress looks :) because there were no failed tasks or missing dependencies

===== Luigi Execution Summary =====
```

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$ hdfs dfs -cat /user/hadoop/output/part-*
2      be
1      the
3      jack
1      over
1      quick
1      jumped
1      nimble
1      candlestick
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Luigi$
```

LAB3: MapReduce

HadoopStreaming

Mapper.py:

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/HadoopStreaming$ hdfs dfs -cat /user/hadoop/input.txt | python2 mapper.py
jack 1
be 1
nimble 1
jack 1
be 1
quick 1
jack 1
jumped 1
over 1
the 1
candlestick 1
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/HadoopStreaming$
```

Reducer.py:


```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/HadoopStreaming$ hdfs dfs -
cat /user/hadoop/input.txt | python2 mapper.py | sort > mapper_output.txt
cat mapper_output.txt | python2 reducer.py
be      2
candlestick      1
jack      3
jumped      1
nimble      1
over      1
quick      1
the      1
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/HadoopStreaming$
```

Mrjob

Word_count.py

```
PROBLEMS 14 OUTPUT DEBUG CONSOLE TERMINAL PORTS 2
"INSPECTIONS,A49102,TRANS-Highways"      2
"INSPECTOR"      199
"INSPECTOR,A06011,Housing"      14
"INSPECTOR,A06013,Housing"      7
"INSPECTOR,A06014,Housing"      9
"INSPECTOR,A06038,Housing"      6
"INSPECTOR,A06043,Housing"      12
"INSPECTOR,A06083,Housing"      11
"INSPECTOR,A14003,FIN-Collections"      5
"INSPECTOR,A50804,DPW-Water"      4
"INSTALLATI,A49506,TRANS-Highways"      1
"INSTALLER"      133
"INSTRUCT,A49102,TRANS-Highways"      1
"INSTRUCTOR,A04009,R&P-Recreation"      4
"INSTRUCTOR,A04010,R&P-Recreation"      1
"INSTRUCTOR,A06076,Housing"      2
"INSTRUCTOR,A65200,HLTH-Health"      2
"INSTRUCTOR,A67002,R&P-Administration"      1
"INSTRUCTOR,A67003,R&P-Administration"      3
"INSTRUCTOR,A67005,R&P-Administration"      1
"INSTRUCTOR,A67011,R&P-Administration"      5
"INSTRUCTOR,A67012,R&P-Administration"      7
"INSTRUCTOR,A68005,R&P-Parks"      17
"INSTRUCTOR,A68011,R&P-Parks"      1
"INSTRUCTOR,B68040,R&P-Parks"      1
"INSTRUCTOR,B68051,R&P-Parks"      2
"INSTRUCTOR,P04001,R&P-Recreation"      252
Removing temp directory /tmp/word_count.hadoop.20250329.163409.891759...
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/mrjob$ python2 word_count.p
y salaries.csv
```

Top_salary.py

```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/mrjob$ python2 top_salary.p
y salaries.csv
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/top_salary.hadoop.20250329.163513.936739
Running step 1 of 1...

Counters: 1
warn
missing gross=3223

job output is in /tmp/top_salary.hadoop.20250329.163513.936739/output
Streaming final output from /tmp/top_salary.hadoop.20250329.163513.936739/output...
"salary"      [163200.0, "\"Chow,Rudolph S\"",DIRECTOR PUBLIC WORKS,A41101,DPW-Administration ,02/01/2011,$163200.00,$145513.
79"]
"salary"      [163200.0, "\"Nilson,George A\"",CITY SOLICITOR,A30001,Law Department ,01/16/2007,$163200.00,$164332.32"]
"salary"      [163365.0, "\"Rawlings-Blake,Stephanie C\"",MAYOR,A01001,Mayors Office ,12/07/1995,$163365.00,$161219.24"]
"salary"      [165000.0, "\"Ford,Niles R\"",EXECUTIVE LEVEL III,A64006,Fire Department ,01/15/2014,$165000.00,$69807.64"]
"salary"      [172000.0, "\"Parthemos,Kaliopel\"",EXECUTIVE LEVEL III,A01020,Mayor's Office ,12/26/2006,$172000.00,$154654.39"]
]
"salary"      [187200.0, "\"Swift,Michael\"",CONTRACT SERV SPEC II,A02003,City Council ,05/19/2008,$187200.00,$3510.00"]
"salary"      [190000.0, "\"Black,Harry E\"",EXECUTIVE LEVEL III,A23001,FIN-Admin & Budgets ,01/30/2012,$190000.00,$188328.50
"]
"salary"      [193800.0, "\"Batts,Anthony W\"",EXECUTIVE LEVEL III,A99390,Police Department ,09/25/2012,$193800.00,$193653.69
"]
"salary"      [200000.0, "\"Charles,Ronnie E\"",EXECUTIVE LEVEL III,A83001,HR-Human Resources ,07/05/2012,$200000.00,$185741.
81"]
```

Avg.py:

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/MapReduce/mrjob$ python2 avg.py salaries.csv
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/avg.hadoop.20250329.163640.264128
Running step 1 of 2...

Counters: 65
  depts
    COMP-Audits=40
    COMP-Communication Ser=30
    COMP-Comptroller's O=12
    COMP-Real Estate=8
    Circuit Court=140
    City Council=93
    Civil Rights & Wage Enfor=12
    Council Services=6
    DPW-Administration=116
    DPW-Solid Waste=255
    DPW-Solid Waste (wkly)=475
    DPW-Water & Waste Water=1491
    ERS/EOS Admin=33
    Elections=8
    Enoch Pratt Free Library=463
    FIN-Acct & Payroll=52
    FIN-Admin & Budgets=31
    FIN-Collections=129
    FIN-Purchasing=52

```

Lab4: Spark

Word_count.py

```

Successfully installed py4j-0.10.9 pyspark-3.1.2
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Spark$ spark-submit --master local[4] word_count.py
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform (file:/opt/spark/jars/spark-unsafe_2.12-3.1.2.jar) to constructor java.nio.DirectByteBuffer(long,int)
WARNING: Please consider reporting this to the maintainers of org.apache.spark.unsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
25/03/30 01:00:54 INFO SparkContext: Running Spark version 3.1.2
25/03/30 01:00:54 INFO ResourceUtils: =====
25/03/30 01:00:54 INFO ResourceUtils: No custom resources configured for spark.driver.
25/03/30 01:00:54 INFO ResourceUtils: =====
25/03/30 01:00:54 INFO SparkContext: Submitted application: SparkWordCount
25/03/30 01:00:54 INFO ResourceProfile: Default ResourceProfile created, executor resources: Map(cores -> name: cores, amount: 1, script: , vendor: , memory -> name: memory, amount: 1024, script: , vendor: , offHeap -> name: offHeap, amount: 0, script: , vendor: ), task resources: Map(cpus -> name: cpus, amount: 1.0)

```

```

hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Spark$ hdfs dfs -cat /user/hduser/output/*
('nimble', 1)
('quick', 1)
('candlestick', 1)
('jack', 3)
('be', 2)
('jumped', 1)
('over', 1)
('the', 1)
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Spark$

```

Text_search.py:


```
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Spark$ spark-submit text_search.py "Toy Story"
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform (file:/opt/spark/jars/spark-unsafe_2.12-3.1.2.jar) to constructor java.nio.DirectByteBuffer(long,int)
WARNING: Please consider reporting this to the maintainers of org.apache.spark.unsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
25/03/30 01:36:25 INFO SparkContext: Running Spark version 3.1.2
25/03/30 01:36:25 INFO ResourceUtils: =====
25/03/30 01:36:25 INFO ResourceUtils: No custom resources configured for spark.driver.
25/03/30 01:36:25 INFO ResourceUtils: =====
25/03/30 01:36:25 INFO SparkContext: Submitted application: SparkWordCount
25/03/30 01:36:25 INFO ResourceProfile: Default ResourceProfile created, executor resources: Map(cores -> name: cores, amount: 1, script: , vendor: , memory -> name: memory, amount: 1024, script: , vendor: , offHeap -> name: offHeap, amount: 0, script: , vendor: ), task resources: Map(cpus -> name: cpus, amount: 1.0)
25/03/30 01:36:25 INFO ResourceProfile: Limiting resource is cpu
25/03/30 01:36:25 INFO ResourceProfileManager: Added ResourceProfile id: 0
```

```
1 Matching titles found:
Toy Story (1995)
25/03/30 01:36:29 INFO SparkUI: Stopped Spark web UI at http://hadoop-2274802010480:4040
25/03/30 01:36:29 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
25/03/30 01:36:29 INFO MemoryStore: MemoryStore cleared
25/03/30 01:36:29 INFO BlockManager: BlockManager stopped
25/03/30 01:36:29 INFO BlockManagerMaster: BlockManagerMaster stopped
25/03/30 01:36:29 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
25/03/30 01:36:29 INFO SparkContext: Successfully stopped SparkContext
25/03/30 01:36:30 INFO ShutdownHookManager: Shutdown hook called
25/03/30 01:36:30 INFO ShutdownHookManager: Deleting directory /tmp/spark-d4a7bd54-d60c-492a-a303-da64fcc9240a/pyspark-4c0bb809-462c-4e34-aa47-a5948db51944
25/03/30 01:36:30 INFO ShutdownHookManager: Deleting directory /tmp/spark-d4a7bd54-d60c-492a-a303-da64fcc9240a
25/03/30 01:36:30 INFO ShutdownHookManager: Deleting directory /tmp/spark-f2d13367-7f18-4396-a11e-b9c022cb9435
hadoop@hadoop-2274802010480:~/iDragonCloud/71ITDS40303/Thuc_hanh/HadoopWithPython/python/Spark$
```

11 4 Zachary Radtka (9 years ago) Ln 28, Col 50 Spaces: 3 UTF-8 LF {} Python 3.12.3 64-bit

```
on/python/Spark/text_search.py:33, look 1735711
1 Matching titles found:
Dead Man Walking (1995)
25/03/30 01:38:01 INFO SparkUI: Stopped Spark web UI at http://hadoop-2274802010480:4040
25/03/30 01:38:01 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
25/03/30 01:38:01 INFO MemoryStore: MemoryStore cleared
25/03/30 01:38:01 INFO BlockManager: BlockManager stopped
25/03/30 01:38:01 INFO BlockManagerMaster: BlockManagerMaster stopped
25/03/30 01:38:01 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
```